

# Snowflake Architecture Demo: The "Sony Pictures" Global Data Platform

Target Audience: Sony Pictures Entertainment (SPE) Data Architecture Team

Objective: Demonstrate a modern, governed Medallion Architecture on Snowflake, highlighting the interoperability of imperative (Streams/Tasks) and declarative (Dynamic Tables) paradigms, integrated with dbt and Data Metric Functions (DMFs).

---

## 1. Demo Architecture Overview (3-10-2 Model)

This demo simulates the processing of global box office data, streaming viewership (Sony Pictures Core), and fan loyalty interactions. We define a precise **3-10-2** table structure.

### The Bronze (Raw) Layer: 3 Tables

*Landing zone for immutable, raw ingestion.*

1. **RAW.SPE.FAN\_INTERACTIONS\_STREAM**: Real-time stream of fan site logins, "Sony Rewards" signups, and movie reviews (JSON format).
2. **RAW.SPE.TITLE\_METADATA\_CATALOG**: Master metadata from internal CMS combined with IMDb/Gracenote feeds (Title, Cast, Director, Genre, Runtime) (JSON format).
3. **RAW.SPE.GLOBAL\_BOX\_OFFICE**: Daily CSV dumps from theatrical distribution partners (AMC, Regal, International Distributors).

### The Silver (Refined) Layer: 10 Tables

*Cleanse, Conform, and Deduplicate.*

- **The Fan Identity Fork (Imperative / Tasks):**
  - Scenario: Handling PII and separating "Verified Purchasers" from "General Web Traffic".
    1. SILVER.STG.STG\_FAN\_VERIFIED: Fans with linked ticket purchases (High Value).
    2. SILVER.STG.STG\_FAN\_GUEST: Anonymous or unverified interactions.
    3. SILVER.DIMS.DIM\_FANS\_SCD2: Historical fan dimension (tracking Region/Preference changes over time).
    4. SILVER.SECURE.DIM\_FAN\_PII\_VAULT: Sensitive data (Email, Phone) isolated in a restricted schema.
  - **The Content & Performance Chain (Declarative / Dynamic Tables):**
    - Scenario: Standardizing Movie Titles and Revenue across regions.
      5. SILVER.STG.STG\_TITLES\_PARSED: Flattened JSON metadata (Cast/Crew arrays).
      6. SILVER.DIMS.DIM\_TITLES\_SCD1: "Golden Record" of movie metadata (Type 1).
      7. SILVER.STG.STG\_BOX\_OFFICE\_FLATTEN: Parsed raw sales files.

8. SILVER.STG.STG\_BOX\_OFFICE\_CURRENCY: Normalizing JPY/GBP/EUR to USD.
9. SILVER.STG.STG\_BOX\_OFFICE\_DEDUP: Deduplication logic using QUALIFY.
10. SILVER.FACTS.FACT\_DAILY\_PERFORMANCE: Final enriched fact table joining Revenue to Titles and Release Dates.

## The Gold (Aggregated) Layer: 2 Tables

*Business-ready aggregations for Tableau/ThoughtSpot.*

1. GOLD.STUDIO\_OPS.AGG\_FRANCHISE\_PERFORMANCE: Dynamic Table (Spider-Verse, Ghostbusters, etc.).
  2. GOLD.MARKETING.AGG\_FAN\_LIFETIME\_VALUE: Dynamic Table (Rolling 365-day spend per fan).
- 

## 2. Deep Dive: The Imperative Path (Streams & Tasks)

**Scenario:** Processing **Fan Loyalty Data**. Because this involves conditional routing (Splitting high-value Verified Fans from casual traffic) and strict PII compliance, we use **Streams and Tasks**.

### Step 2.1: The Split (Conditional Multi-Table Insert)

We use a Task to read from a Stream and "fan out" the data.

SQL

```
-- 1. Create the Stream on the raw JSON landing table
CREATE STREAM raw_fan_stream ON TABLE raw.spe.fan_interactions_stream;

-- 2. Create the Router Task
CREATE OR REPLACE TASK tsk_route_fan_data
WAREHOUSE = 'transform_wh'
SCHEDULE = '1 minute'
WHEN SYSTEM$STREAM_HAS_DATA('raw_fan_stream')
AS
-- "The Split": Conditional Multi-Table Insert
INSERT ALL
-- Path A: Verified Ticket Buyers (Linked to Fandango/Cinema account)
WHEN raw_json:account_type::STRING = 'VERIFIED_LINKED' THEN
    INTO silver.stg.stg_fan_verified
```

```

VALUES (fan_id, email, region, 'VERIFIED', event_ts)
-- Path B: Guest/Anonymous Traffic
ELSE
    INTO silver.stg.stg_fan_guest
    VALUES (session_id, ip_address, region, 'GUEST', event_ts)
SELECT
    raw_json:id::STRING as fan_id,
    raw_json:session::STRING as session_id,
    raw_json:email::STRING as email,
    raw_json:ip::STRING as ip_address,
    raw_json:region::STRING as region,
    raw_json:event_time::TIMESTAMP as event_ts,
    raw_json
FROM raw_fan_stream;

```

## Step 2.2: The Finalizer (SCD Type 2 Merge)

We chain a second task to process the Verified Fans into the Dimension, preserving history (e.g., A fan moves from "UK" to "USA" - we want to track that migration).

SQL

```

CREATE OR REPLACE TASK tsk_merge_fan_scd2
AFTER tsk_route_fan_data
AS
MERGE INTO silver.dims.dim_fans_scd2 AS target
USING silver.stg.stg_fan_verified AS source
ON target.fan_id = source.fan_id
-- Complex Logic to close old record (set valid_to) and insert new one
... ;

```

---

## 3. Deep Dive: The Declarative Path (Dynamic Tables)

**Scenario:** Processing **Box Office & Title Data**. This is a linear analytical flow (Ingest -> Clean -> Currency Conversion -> Aggregate). We use **Dynamic Tables (DTs)**.

### The "Cleanse, Dedup, Lookup" Chain

### Step 3.1: Cleanse & Flatten (Table 7)

SQL

```
CREATE DYNAMIC TABLE silver.stg.stg_box_office_flatten
TARGET_LAG = '5 minutes'
WAREHOUSE = 'transform_wh'
AS
SELECT
    raw_data:theater_id::STRING as theater_id,
    raw_data:title_id::STRING as title_id,
    raw_data:local_gross::FLOAT as local_gross_amt,
    raw_data:currency::STRING as currency_code,
    raw_data:show_date::DATE as report_date
FROM raw.spe.global_box_office;
```

### Step 3.2: Currency Normalization & Data Quality (Table 8)

*Note: We attach a DMF here later.*

SQL

```
CREATE DYNAMIC TABLE silver.stg.stg_box_office_currency
TARGET_LAG = '5 minutes'
AS
SELECT
    b.title_id,
    b.report_date,
    -- Simple conversion logic (Real-world would join to a Rate table)
    CASE
        WHEN b.currency_code = 'JPY' THEN b.local_gross_amt * 0.0067
        WHEN b.currency_code = 'GBP' THEN b.local_gross_amt * 1.27
        ELSE b.local_gross_amt
    END AS gross_usd
FROM silver.stg.stg_box_office_flatten b
WHERE gross_usd >= 0; -- Basic filtering
```

### Step 3.3: Deduplicate (Table 9)

The critical "Dedup" step using window functions.

SQL

```
CREATE DYNAMIC TABLE silver.stg.stg_box_office_dedup
TARGET_LAG = '5 minutes'
AS
SELECT *
FROM silver.stg.stg_box_office_currency
QUALIFY ROW_NUMBER() OVER (
    PARTITION BY title_id, report_date, theater_id
    ORDER BY ingestion_ts DESC
) = 1;
```

- **Why:** QUALIFY ensures that if a theater sends a correction file for the same date, we only process the latest version.<sup>1</sup>

### Step 3.4: Dimension Lookup & Fact Creation (Table 10)

Joining the sales stream with the Movie Dimension to create the Star Schema.

SQL

```
CREATE DYNAMIC TABLE silver.facts.fact_daily_performance
TARGET_LAG = '10 minutes'
AS
SELECT
    s.report_date,
    s.gross_usd,
    t.title_key, -- Surrogate Key from Titles Dim
    t.franchise_name -- e.g., 'Spider-Man Universe'
FROM silver.stg.stg_box_office_dedup s
LEFT JOIN silver.dims.dim_titles_scd1 t
    ON s.title_id = t.imdb_id;
```

## 4. Governance: Data Metric Functions (DMFs)

We integrate **Data Quality** directly into the pipeline.

### Step 1: Define the Metric (Valid IMDb Score)

SQL

```
CREATE DATA METRIC FUNCTION governance.metrics.valid_imdb_score(  
    ARG_T FLOAT  
)  
RETURNS NUMBER AS  
$$ SELECT CASE WHEN ARG_T BETWEEN 0 AND 10 THEN 1 ELSE 0 END$$;
```

### Step 2: Apply to Title Metadata

We apply this to DIM\_TITLE\_SCD1 to ensure we don't ingest corrupted rating data.

SQL

```
ALTER TABLE silver.dims.dim_titles_scd1  
ADD DATA METRIC FUNCTION governance.metrics.valid_imdb_score(imdb_rating)  
ON SCHEDULE '5 MINUTE';
```

### Step 3: View Results

Any failures are automatically logged to  
SNOWFLAKE.LOCAL.DATA\_QUALITY\_MONITORING\_RESULTS.

---

## 5. dbt Integration

Why dbt for Sony Pictures?

While Snowflake handles the execution, dbt handles the engineering lifecycle.

1. **Macros for Currency Conversion:** Instead of hardcoding exchange rates in SQL (as seen in Step 3.2), we use a dbt macro {{ convert\_currency('amount', 'currency\_code') }} ensuring all analysts use the same official Finance logic.

2. **Environment awareness:** Run the pipeline in DEV\_SPE\_DB during testing and PROD\_SPE\_DB for release seamlessly.
  3. **Docs:** dbt generates a documentation site showing that FACT\_DAILY\_PERFORMANCE is downstream of DIM\_TITLES, helping Business Analysts understand the data lineage.<sup>2</sup>
- 

## 6. Gold Layer Analytics: 10 Example Queries

This section demonstrates the value of the Gold Layer to Sony Pictures executives. We categorize these by complexity.

### Simple Queries (The "What Happened?" Layer)

#### 1. Franchise Leaderboard (Box Office)

- *Objective:* Which Franchise is generating the most revenue this quarter?

SQL

```
SELECT
    t.franchise_name,
    SUM(f.gross_usd) as total_global_box_office
FROM gold.studio_ops.agg_franchise_performance f
JOIN silver.dims.dim_titles_scd1 t ON f.title_key = t.title_key
WHERE f.report_date >= DATE_TRUNC('QUARTER', CURRENT_DATE())
GROUP BY 1
ORDER BY 2 DESC;
```

#### 2. Top Rated Action Movies

- *Objective:* Identify high-quality back-catalogue content for Sony Pictures Core promotion.

SQL

```
SELECT
    title_name,
    release_year,
    imdb_rating
```

```
FROM silver.dims.dim_titles_scd1
WHERE genre = 'Action' AND release_year < 2020
ORDER BY imdb_rating DESC
LIMIT 10;
```

### 3. Fan Regional Distribution

- *Objective:* Where are our "Super Fans" located?

SQL

```
SELECT
    region,
    COUNT(DISTINCT fan_id) as active_verified_fans
FROM silver.dims.dim_fans_scd2
WHERE is_current = TRUE
    AND loyalty_tier = 'PLATINUM'
GROUP BY 1;
```

### 4. Daily Release Tracking

- *Objective:* How is a specific movie performing day-over-day since release?

SQL

```
SELECT
    report_date,
    gross_usd
FROM silver.facts.fact_daily_performance
WHERE title_key = (SELECT title_key FROM silver.dims.dim_titles_scd1 WHERE title_name =
    'Kraven the Hunter')
ORDER BY report_date ASC;
```

## Medium Queries (The "Why/Correlation" Layer)

### 5. ROI Analysis: Marketing vs. Box Office

- *Objective:* Compare marketing spend (ingested from SAP) against opening weekend

actuals.

SQL

```
SELECT
    t.title_name,
    m.marketing_budget,
    SUM(f.gross_usd) as opening_weekend_gross,
    (SUM(f.gross_usd) / m.marketing_budget) * 100 as roi_percentage
FROM silver.facts.fact_daily_performance f
JOIN silver.dims.dim_titles_scd1 t ON f.title_key = t.title_key
JOIN raw.spe.marketing_spend m ON t.title_id = m.title_id
WHERE f.report_date BETWEEN t.release_date AND t.release_date + 2
GROUP BY 1, 2;
```

## 6. Fan Affinity Analysis (Basket Analysis)

- Objective: "If a fan watched *Spider-Man*, what else did they watch?"

SQL

```
WITH spidey_fans AS (
    SELECT DISTINCT fan_key
    FROM silver.facts.fact_watch_events
    WHERE title_key = (SELECT title_key FROM silver.dims.dim_titles_scd1 WHERE title_name =
'Spider-Man: Across the Spider-Verse')
)
SELECT
    t.title_name,
    COUNT(DISTINCT w.fan_key) as shared_viewers
FROM silver.facts.fact_watch_events w
JOIN silver.dims.dim_titles_scd1 t ON w.title_key = t.title_key
WHERE w.fan_key IN (SELECT fan_key FROM spidey_fans)
    AND t.title_name!= 'Spider-Man: Across the Spider-Verse'
GROUP BY 1
ORDER BY 2 DESC
LIMIT 5;
```

## 7. Sentiment Trends for Directors

- *Objective:* Analyze how fan sentiment changes for a director over their last 3 films.

SQL

```
SELECT
    t.director,
    t.title_name,
    t.release_year,
    AVG(f.sentiment_score) as avg_fan_sentiment
FROM silver.facts.fact_fan_reviews f
JOIN silver.dims.dim_titles_scd1 t ON f.title_key = t.title_key
WHERE t.director = 'Christopher Nolan' -- Example
GROUP BY 1, 2, 3
ORDER BY 3 ASC;
```

## 8. SCD2 Analysis: The "Moved Fan" Effect

- *Objective:* Do fans spend more or less after moving regions? (Leveraging SCD2 history).

SQL

```
SELECT
    curr.region as new_region,
    prev.region as old_region,
    AVG(curr.annual_spend - prev.annual_spend) as spend_delta
FROM silver.dims.dim_fans_scd2 curr
JOIN silver.dims.dim_fans_scd2 prev
    ON curr.fan_id = prev.fan_id
    AND curr.previous_valid_to = prev.valid_to -- Linking history chains
WHERE curr.region!= prev.region
GROUP BY 1, 2;
```

## Complex Queries (The "Predictive/Strategic" Layer)

### 9. The "Blockbuster Prediction" Model

- *Objective:* Correlate pre-release social sentiment (Bronze Stream) with Post-Release Box Office (Silver Fact) to create a predictive multiplier for future films.

SQL

```

WITH PreReleaseBuzz AS (
    -- Aggregate social sentiment 30 days BEFORE release
    SELECT
        title_id,
        COUNT(*) as mention_volume,
        AVG(sentiment_score) as buzz_score
    FROM raw.spe.fan_interactions_stream
    WHERE event_type = 'SOCIAL_MENTION'
    GROUP BY 1
),
FirstWeekPerformance AS (
    -- Aggregate Revenue 7 days AFTER release
    SELECT
        t.imdb_id as title_id,
        SUM(f.gross_usd) as week1_revenue
    FROM silver.facts.fact_daily_performance f
    JOIN silver.dims.dim_titles_scd1 t ON f.title_key = t.title_key
    WHERE f.report_date <= t.release_date + 7
    GROUP BY 1
)
SELECT
    t.genre,
    CORR(p.buzz_score, w.week1_revenue) as correlation_coefficient,
    -- If buzz is high but revenue is low, we have a marketing misalignment
    AVG(CASE WHEN p.buzz_score > 0.8 THEN w.week1_revenue ELSE NULL END) as
high_buzz_avg_rev
FROM PreReleaseBuzz p
JOIN FirstWeekPerformance w ON p.title_id = w.title_id
JOIN silver.dims.dim_titles_scd1 t ON p.title_id = t.imdb_id
GROUP BY 1;

```

## 10. "Super Fan" Lifetime Value (LTV) Cohort Analysis

- *Objective:* Calculate LTV across Box Office, Streaming, and Merch, grouped by the year they joined the loyalty program.

SQL

```
SELECT
    DATE_TRUNC('YEAR', d.join_date) as cohort_year,
    COUNT(DISTINCT d.fan_key) as cohort_size,
    -- Aggregate spend from multiple fact tables
    SUM(COALESCE(bo.ticket_spend, 0) + COALESCE(str.sub_spend, 0) +
    COALESCE(mer.merch_spend, 0)) as total_cohort_revenue,
    (total_cohort_revenue / cohort_size) as avg_ltv
FROM silver.dims.dim_fans_scd2 d
-- Left Join to Box Office Spend
LEFT JOIN (
    SELECT fan_key, SUM(amount) as ticket_spend FROM silver.facts.fact_ticket_sales GROUP BY 1
) bo ON d.fan_key = bo.fan_key
-- Left Join to Streaming Spend
LEFT JOIN (
    SELECT fan_key, SUM(amount) as sub_spend FROM silver.facts.fact_streaming_subs GROUP
    BY 1
) str ON d.fan_key = str.fan_key
-- Left Join to Merchandise Spend
LEFT JOIN (
    SELECT fan_key, SUM(amount) as merch_spend FROM silver.facts.fact_merch_sales GROUP BY
    1
) mer ON d.fan_key = mer.fan_key
WHERE d.is_current = TRUE
GROUP BY 1
ORDER BY 1 DESC;
```

## 7. Summary: Decision Matrix for SPE

| Feature              | Dynamic Tables<br>(Declarative)  | Streams & Tasks<br>(Imperative)   |
|----------------------|--|---|
| <b>Best Used For</b> | Continuous pipelines<br>(Global Box Office \$\\to\$ Silver \$\\to\$ Gold). "I want | Complex orchestration,<br>routing, or non-SQL logic.<br>"Check if Fan is Verified; if |

|                     |   |   |
|---------------------|---|---|
|                     | the daily revenue numbers, keep them fresh."  | yes, PII Vault; if no, Guest Table."  |
| <b>Data Logic</b>   | <b>Linear Chaining:</b> Great for currency conversion, parsing Cast/Crew JSON arrays, and aggregating franchise totals. | <b>Splitting/Routing:</b> INSERT ALL allows splitting the Fan Stream into multiple security zones (Public vs. PII). |
| <b>SCD Strategy</b> | <b>SCD Type 1:</b> Auto-updates Movie Metadata (e.g., Runtime correction).  | <b>SCD Type 2:</b> Tracking Fan migration (Region changes) for historical compliance.                               |
| <b>Maintenance</b>  | <b>Low:</b> Snowflake manages the scheduler.  | <b>Medium:</b> You manage the DAG dependencies.   |
| <b>Our Demo</b>     | Used for <b>Content &amp; Revenue</b> (High volume, linear transformation).   | Used for <b>Fan Loyalty &amp; PII</b> (Complex validation, routing, History tracking).                              |

## Works cited

1. Deduplicating CDC records in Snowflake: The fastest way | by Emiliano Mancuso - Medium, accessed January 21, 2026,  
<https://medium.com/@emancu/deduplicating-cdc-records-in-snowflake-the-fastest-way-5a4a14f9890a>
2. Accelerating Data Teams with dbt Cloud & Snowflake, accessed January 21, 2026,  
<https://www.snowflake.com/en/developers/guides/data-teams-with-dbt-cloud/>